



Mustikarasa System Architecture Overview (TOGAF-Based Approach)

Introduction and Scope

The **Mustikarasa Project** is an effort to turn a historical Indonesian cookbook, *Mustika Rasa*, into a reliable digital edition that serves multiple purposes. This document presents a comprehensive architecture overview of the Mustikarasa system, treating it from an enterprise architecture perspective (inspired by TOGAF's domains: Business, Application, Data, and Technology). The goal is to provide a “total picture” of the system – covering business and technical requirements, the system’s capabilities and their maturity, and the future roadmap (“architectural runway”) – in alignment with the project’s vision and strategy. The architecture is described in broad terms, emphasizing traceability, agent-based workflows, and governance, as per the guiding principles of the project [1](#) [2](#).

Context: The Mustikarasa system is not merely a translation pipeline; it is an **agentic workflow** that uses multiple AI agents orchestrated to simulate an editorial process with human-like rigor. Each piece of content (a recipe or section of the book) passes through a series of specialized agents (for translation quality, readability editing, cultural annotation, etc.), coordinated by a central Orchestrator agent [3](#) [4](#). Throughout this process, every change is logged and justified, ensuring the system can always show “*zo was het — zo is het — en dit is waarom het veranderde*” (“this is how it was, this is how it is, and this is why it changed”) [2](#). This overarching principle of traceability and accountability is the cornerstone of both the business requirements and the technical design.

In the following sections, we detail the **Business Architecture** (the goals, requirements, and process), the **Application Architecture** (the components – agents and their interactions), the **Data/Information Architecture** (how content and metadata are managed), and the **Technical Architecture** (technology stack and infrastructure). We then list the system’s **Capabilities and their Maturity**, and finally provide **Recommendations** to guide future development. This architecture baseline implicitly reflects the project’s vision and strategy [5](#), ensuring that the system design supports the long-term mission of creating a trustworthy, traceable, and reusable knowledge infrastructure [6](#).

Business Context and Requirements

The Mustikarasa project exists to **restore and modernize a historical text** in a trustworthy way. Several key business drivers and requirements shape the system’s architecture:

- **Authenticity and Integrity:** The source material (*Mustika Rasa*) is a historic document, and the project’s aim is to **preserve its integrity** while making it accessible [7](#). Digital versions of this text have been incomplete or error-prone historically (due to OCR errors and scan issues), so the system must produce a reliable “*digital master*” that reflects the original book as faithfully as possible [8](#). All technical corrections (OCR fixes, structural repairs) should be made **without altering the historical**

content, except to fix obvious transcription errors ⁹. In other words, *technical errors are corrected, but content errors are retained and documented* ¹⁰ – the system should **never rewrite history**, only **document it**.

- **Traceable Editorial Process:** Rather than simply creating a new edition, the project prioritizes a **traceable process**. Every change or decision in editing must be documented with its rationale. The core business principle is that the system can show for any element: “*here is how it was originally, here is how it is now, and here is why it was changed*” ². This means **all edits must be auditable**: the original state, the modified state, and the reason for each modification are recorded ¹¹. This requirement ensures **accountability** in the editorial workflow and builds trust in the final product.
- **Multi-Phase Mission:** The project is divided into **three sequential missions (phases)**, each with distinct goals that translate into system requirements ¹² ¹³:
 - **Phase 0 – Source Restoration:** Create a *reliable digital master* of the original book ⁸. The system must support tasks like OCR correction, reconstruction of tables and structure, image linking, and detection of missing or damaged content ⁹. All restoration actions must be reversible and logged; for example, if a typo in OCR is fixed, the system logs the original text, the corrected text, and the reason for change ¹¹.
 - **Phase 1 – Scholarly Edition:** Produce a *scholarly annotated edition* that remains faithful to the original text ¹³. In this phase, **content is not “improved” or modernized**; instead, errors or contradictions in the original are preserved but explicitly flagged and annotated ¹⁴. The system needs to differentiate between the pure source text and the annotated version ¹⁵, supporting critical commentary (footnotes, annotations) without altering the source content. Business-wise, this serves researchers and editors who want to see the original content *with scholarly insights*.
 - **Phase 2 – Public Edition:** Deliver a *modernized public edition* that is accessible and engaging to general readers ¹⁶. Here, the system may allow rewriting or correcting recipes for clarity, safety, and readability, **but with full accountability to the scholarly edition** ¹⁷. This means any modernization (e.g. updating obsolete terms, adjusting for safety/allergies) must be traceable back to what it was and why it changed ¹⁸. The public edition should be “*veilig, leesbaar, inspirerend*” (safe, readable, inspiring) ¹⁶ while retaining links to the historical context. This phase prioritizes **usability and design** (narrative flow, visual appeal) in addition to accuracy.
- **Multi-Audience Support:** The system must ultimately serve **multiple audiences** ¹⁹, including academic researchers, project editors, and general readers. This implies that the outputs of the system might be delivered in different formats or levels of detail:
 - For researchers: a rigorously annotated edition (Phase 1 output) with complete traceability and commentary.
 - For general readers: a refined edition (Phase 2 output) that reads smoothly in modern language, with critical changes justified perhaps in an appendix or interactive notes.
 - Internally, for project editors/governance: detailed logs and decision records to audit the process. The architecture must accommodate these different end-products from a single source, ensuring consistency between them (e.g., any content change in the public edition is documented in terms of the scholarly edition evidence ²⁰).

- **Quality over Speed:** A key strategic choice is that **the project values a correct, repeatable process over rapid output**²¹. The goal is to build a “knowledge infrastructure, not just a one-time translation”²². As a result, business requirements emphasize **rigor, repeatability, and safety**. The system should prevent uncontrolled changes and ensure that if something isn’t verified, it doesn’t silently proceed. (For example, the project has instituted a governance process where any critical uncertainties trigger a stop for review²³.) This means the workflow might be slower due to human checks and iterative pilots, but it ensures the final product is trustworthy. This requirement guides many technical safeguards described later (such as stop mechanisms and human gate reviews).
- **Governance and Compliance:** The business stakeholders require that the system operate under strict **governance rules**. This includes compliance with an internal “Human Gate” process for risky decisions, version control of requirements, and audit trails for decisions²⁴⁵. For instance, the project has a formal **requirements change process** to control how requirements evolve and ensure traceability of requirement changes²⁴. Similarly, any terminological decisions (like how to translate a culinary term) must follow a defined **Glossary Decision Lifecycle** that involves human approval, rather than being decided solely by the AI²⁵. These governance needs translate into system features like logging every decision, requiring manual approval for glossary updates, and halting on incidents – all of which are detailed in the technical requirements.

Below is a summary of **key business requirements** distilled from the above context (and the project’s Vision document), which the architecture must satisfy:

- **BR1. Create an Authentic Digital Master:** The system shall produce a digital version of *Mustika Rasa* that is as faithful as possible to the original printed book (fixing only technical issues like OCR errors)⁸²⁶.
- **BR2. Produce Sequential Editions with Traceability:** The system shall support producing a scholarly annotated edition and a modern public edition from that master, maintaining traceability of every change between editions (original → scholarly → public)¹³²⁷.
- **BR3. Ensure Complete Traceability and Rationale:** For any correction, annotation, or modernization, the system must capture *what was changed, why, and by whom/what*, in a way that can be reviewed later²²⁸. This includes preserving original content alongside changes and logging decision rationales.
- **BR4. Preserve Content Integrity (No Unapproved Changes):** The system shall not allow irreversible or content-changing decisions to be made autonomously by AI agents. All substantive content decisions (e.g. glossary term standardization, removal of content, changes affecting meaning) require human review (the “Human Gate”) before being finalized²⁹²⁵.
- **BR5. Support Multi-Audience Outputs:** The system’s outputs must cater to different user groups – e.g., a detailed academic output and a simplified public output – without duplication of effort or divergence in content integrity¹⁹¹⁶. The architecture should enable reusing the same content base for multiple publication formats.
- **BR6. Maintain High Quality through Governance:** The process must include quality control checkpoints (governance stops, reviews, red-team evaluations) to ensure that no significant errors or unresolved issues propagate to the final output³⁰³¹. “Done” in this context means all known issues are either resolved or explicitly acknowledged and logged – not that the text is 100% error-free³¹.
- **BR7. Build a Reusable Knowledge Infrastructure:** The outcome of the project (data, processes, tools) should be reusable for future research or similar projects⁶. This implies the architecture

should favor open data formats, clear documentation, and modular design so the system can be extended or adapted beyond this specific book.

These business requirements form the basis for the technical architecture decisions. Next, we translate these into specific technical requirements and design principles that the system must follow.

Technical Requirements and Design Principles

To fulfill the above business needs, the Mustikarasa system's design is governed by a set of **technical requirements and principles**. These ensure the system behaves safely and predictably while leveraging advanced AI components. Key technical requirements include:

- **Complete Traceability and Reversibility (TR1):** Every operation that modifies content must be **traceable and reversible** ¹¹. In practice, this means the system keeps a detailed log of changes with before/after states and reasons. For example, if an OCR correction is applied or a sentence is rephrased for readability, the original text and the new text are both stored, along with a justification ²⁶. Additionally, the system must be able to **roll back** changes if needed. A requirement was explicitly added to ensure *safe, reversible batch rollback* of edits, so that if a large batch of changes has issues, it can be undone without losing already approved items ³². This protects against large-scale accidental errors and aligns with the project's "no silent errors" policy.
- **Audit Logging and Decision Recording (TR2):** The system shall **log every step and decision** taken during the workflow ³³. This includes agent outputs, any overrides or interventions, and final outcomes. A specialized *Methodology Archivist* component is planned to compile a **METHODOLOGY_LOG** capturing context, decisions, rationales, and outcomes for each major action ³⁴. Even if that agent is not fully implemented yet, the Orchestrator is already required to send important decisions to be logged ³⁵. The logging must be structured and persistent (e.g., written to a file or database that serves as an audit trail). This ensures that at any point, a reviewer can reconstruct **why the system did what it did**, satisfying traceability requirements. In the Definition of Done, one criterion is that "*all important decisions are traceable*" and "*risks are named and not silently ignored*" ³¹ – robust logging is how the system enforces that.
- **Human-in-the-Loop for Critical Decisions (TR3):** No AI agent is allowed to make final, irreversible content decisions without human approval. In line with the philosophy "agents as signalers, not deciders" ²⁹, the system must implement **gating mechanisms** for certain actions:
- **Glossary / Terminology changes:** If an AI suggests a change in terminology or a translation of a term, that suggestion is treated as a *proposal* only. A formal requirement (ASR-005) was adopted prohibiting automatic glossary updates without a Human Gate review and a rollback plan ³⁶. In practice, the system should route terminological proposals to a human or governance board for approval, and only incorporate them into the official glossary after approval.
- **Incident handling:** If any **incident** or anomaly is detected (e.g., conflicting agent outputs, a potential content risk), the system must *halt further automated processing* and escalate the issue. A requirement (ASR-006) mandates a **governance-stop whenever an incident is raised**, preventing the pipeline from continuing blindly after critical failures ²³. This typically involves invoking a *Troubleshooting Agent* to generate an **INCIDENT_REPORT** and then requiring human intervention to decide how to proceed ³⁷.

- **Publication decisions:** The system cannot autonomously publish or mark content as final without passing human review. For example, the Orchestrator is not allowed to make *definitive publication decisions* on its own ³⁸ – final sign-off is reserved for human decision-makers.

The implementation of these human-in-the-loop controls is via a “**stop model**” with multiple levels of stops ³⁹ ⁴⁰ : - *Soft-stop*: the system pauses to attempt self-correction or seek more information (e.g., ask the Research Agent for clarification) without human involvement ⁴¹ . - *Governance-stop*: the system stops and triggers internal governance agents (e.g., Troubleshooting Agent for an incident report, Methodology Archivist for logging) but still does not involve a human yet ⁴² . - *Hard-stop*: the system halts and requires a human decision (Human Gate) to continue ⁴⁰ . This happens if governance rules conflict or a risk is too high.

This tiered stopping mechanism ensures that minor issues can be resolved in-system, while major ones get human oversight, balancing autonomy with safety.

- **Template and Structure Conformance (TR4):** All agent outputs must conform to predefined templates and structures to ensure consistency. A **Template Agent** provides a contract (a **TEMPLATE_DEFINITION** specifying sections and format) for each type of task ⁴³ . The Orchestrator is required to consult the Template Agent first for each content piece and enforce that subsequent agents' outputs fit the expected structure ⁴³ . If an agent deviates from the template, the Orchestrator will retry or invoke troubleshooting ⁴⁴ . This requirement prevents format drift and makes sure, for example, that every recipe output includes all required sections (ingredients, instructions, annotations, etc.) in a consistent manner. It externalizes the document structure as an explicit artifact, rather than burying it in code, which aligns with governance (structure can be changed via Template Agent with versioning rather than code changes).
- **Reference Integrity and Evidence (TR5):** When generating annotations or making decisions, agents must use authoritative sources and **not hallucinate facts**. The system includes a *Research Agent* that can retrieve a **RESEARCH_REPORT** with relevant references from a curated corpus ⁴⁵ . Technical design mandates that agents like the Annotation Agent or Glossary Agent use the *RelevantFiles* and *KeyInsights* provided by the Research Agent as input context ⁴⁶ . Furthermore, any claim or change that needs evidence (e.g., altering a cooking term for safety reasons) should be linked to source evidence (like original text or an external reference). This principle is part of keeping the system's outputs verifiable. The Orchestrator is **forbidden from searching files or knowledge on its own** as if it were truth ⁴⁷ – it must go through the Research Agent or a human for information, ensuring that all external knowledge is vetted and logged. This design choice enforces that *historical context or prior translations are explicitly consulted and cited*, rather than implicitly assumed by the AI.
- **Consistency and Continuity (TR6):** The system should ensure **consistency across the entire text and process**. This implies two things:

- *Internal consistency of content:* If multiple recipes mention the same term or concept, they should be handled in a consistent way (terminology, translation style, etc.). An agent role (Continuity/Cohesion Agent) is dedicated to checking coherence across sections ⁴⁸ ⁴⁹ . The technical requirement is that before finalizing content, the system runs continuity checks to catch any contradictions or uneven treatments.

- **Process consistency:** The system's behavior must be consistent across runs and over time. This is achieved by having “**golden documents**” for governance – e.g., canonical prompt files, the glossary decision policy, etc., which the Orchestrator must respect ⁵⁰ – and by not embedding hidden logic in code. All agent instructions are centralized in prompt files (the **prompt-as-truth principle**), and the code should not have hard-coded prompts or roles ³⁸. This approach (document-driven logic) ensures that if the process or format changes, it's done through version-controlled documents and templates, maintaining consistency and traceability of the process itself.
- **Robust Error Handling and Isolation (TR7):** Given the complexity of multi-agent interaction, the system must be **robust to errors**. If one part of the workflow fails or produces an inconsistency, the architecture should contain the failure:
 - The Orchestrator is responsible for detecting issues like conflicting agent outputs or template violations and handling them gracefully (retrying or escalating) ⁵¹.
 - The system should isolate failures **per content item** where possible. For example, if processing one recipe in a batch encounters a critical error, that recipe's process might halt (and mark for human review) but this should not corrupt or automatically halt the processing of unrelated recipes. The governance design explicitly tested a *partial batch failure* scenario to ensure that one issue triggers a soft-stop for that item but does not cause unnecessary full batch termination ⁵². The requirement is that batch workflows and the governance model work together such that we avoid cascading failures – only affected items are stopped, and even then, with proper logging and no silent continuation ⁵².
- **No silent corrections:** Any error correction that the system performs (even minor ones) must be logged; nothing should be “fixed” under the hood without leaving a trace ⁵³. If an error is detected but not resolved, it should be explicitly marked and not just ignored. This is part of the Definition of Done: a task is not done unless all errors are either fixed or documented ³¹.
- **Performance and Scalability (TR8):** Although not the top priority, the system should be designed to eventually handle the entire book (hundreds of recipes) in a reasonable time frame. This means the architecture should support **batch processing** and parallelism where safe. The introduction of a *Batch Governor* agent (conceptual) is aimed at managing workflows at the batch (chapter or multiple recipes) level, coordinating parallel runs and ensuring system resources are used efficiently. The technical approach uses CLI scripts and possibly containerized runs (there are references to sandbox configurations for running agents with specific models ⁵⁴). The architecture should allow swapping out or scaling the AI models (e.g., running local models or calling cloud APIs) without changing the core logic – this is facilitated by the *Technical Advisor* agent which can recommend model or configuration changes for performance or cost reasons ⁵⁵. In essence, the design is modular enough that improvements in model performance or computing infrastructure can be slotted in under guidance, to scale up throughput when needed.
- **Security and Safety (TR9):** Any transformations, especially for the public edition, must consider safety (e.g., food safety in recipes) and not introduce liability. The system should flag any content that might be unsafe or require modern disclaimers. For instance, the team added a rule that *health-related statements in the original are preserved and annotated, not casually rewritten*, to avoid changing historical claims ⁵⁶. While this is partly a content policy, it has architectural implications: an agent (possibly the Challenger or a Safety reviewer) should detect such cases and ensure they are handled

via annotation rather than alteration. Additionally, from a security standpoint, if using external AI services or models, the system should not expose sensitive data. In our case, the content is not personal data, but there is a need to protect the integrity of the historical text (no unauthorized modifications). The strict governance stops and human oversight serve as the primary safety net here.

These technical requirements enforce that the system operates under **strict governance, transparency, and control**, in line with the business objectives. Many of these requirements have been codified in project documents and even enforced through development policies (for example, changes to requirements must follow a template and be logged ²⁴). Together, they shape an architecture that is cautious but robust – capable of leveraging AI agents for efficiency while ensuring humans remain in control of the narrative and quality.

Architecture Overview

At a high level, the Mustikarasa system architecture can be viewed as a **pipeline of coordinated AI agents within a governance framework**. The design reflects a layered approach:

- **Business Layer (Process):** The top layer is the editorial process modeled in phases (source restoration, scholarly annotation, public edition). This defines *what* needs to happen in each stage from a content perspective, as described in the business context. Each phase's tasks are supported by specific agents or modules in the application layer below.
- **Application Layer (Agents & Workflow):** The core of the system is an **Orchestrator** and a collection of specialized **Worker Agents**, forming an orchestrator-worker multi-agent architecture. The Orchestrator acts as a **manager** that drives each piece of content through the required workflow, calling on various agents in sequence or in loops as needed ³ ⁵⁷. Each **Agent** is an AI component (powered by Large Language Models) with a well-defined role and output format. For example:
 - The *Translation Quality Agent* ensures the initial translation of a recipe is accurate in meaning.
 - The *Readability Editor* refines the text to be fluent and easy to read without altering meaning.
 - The *Cultural-Historical Agent* (also called Cultural-Historical Editor) adds context or annotations about cultural or historical aspects of the recipe.
 - The *Annotation Agent* inserts scholarly annotations/footnotes highlighting errors or important context.
 - The *Recipe, Table, Image Agents* handle domain-specific content (formatting recipe steps, reconstructing tables, attaching images).
 - The *Fidelity Agent* compares the final edited text with the original (or English reference) to note any meaning deviations ⁵⁸ ⁵⁹.
 - The *Continuity/Cohesion Agent* checks consistency across the content (terminology, style continuity across chapters).
 - The *Design Agent* might apply stylistic guidelines or ensure the output is formatted attractively (especially for the public edition).
 - The *Challenger Agent* provides a form of “devil’s advocate” review – it might re-evaluate decisions or offer alternative suggestions to ensure the best outcome (for instance, catching issues others missed).

These agents correspond to the capabilities the project identified as necessary for a comprehensive editorial workflow. Many of them are implemented as of now (active in code) and are invoked by the Orchestrator during a run ⁶⁰ ⁴⁹. The Orchestrator ensures that **each content item (e.g., a recipe)** passes through all relevant agents in the correct order, and that their outputs are compiled into a coherent final result ³ ⁴. It's akin to an assembly line where each agent is a station adding or checking something.

What makes this more than a simple pipeline is the **governance and feedback loop**: the Orchestrator doesn't blindly move from agent to agent; at each step it monitors compliance and outcomes. If, for example, an agent's output is missing a required section or two agents' suggestions conflict, the Orchestrator can pause and either retry a step or call a specialized agent for help (like Troubleshooting) ⁵¹. The Orchestrator is programmed to **log every step** ³³ and to enforce policies such as "don't accept glossary changes without approval" (it will treat Glossary Agent output as advisory only) ²⁵. In effect, the Orchestrator embodies the rules set in the technical requirements section, acting as the first line of governance control in the live workflow.

- **Governance Layer (Oversight & Logging):** Parallel to the content processing agents, the architecture includes governance-focused agents that oversee and record the process:
- The *Template Agent* and *Research Agent* can be thought of as **supporting agents** that provide structure and context at the Orchestrator's request. The Template Agent gives the format for outputs ⁴³, and the Research Agent provides historical context and references from the document repository ⁴⁵. They don't finalize content; they guide other agents.
- The *Glossary Agent* produces terminology change proposals if needed, ensuring consistent terminology usage, but (per policy) its output goes through human validation before adoption ²⁵.
- The *Methodology Archivist* is an agent tasked with logging the decision process (writing the METHODOLOGY_LOG) whenever key events occur ³⁴. For example, if a non-standard process is executed or an exception is handled, the Orchestrator will invoke the Methodology Archivist to document what happened and why ⁶¹. This agent ensures the **audit trail** is compiled in a human-readable form.
- The *Technical Strategy Advisor* (Technical Advisor) can be invoked to provide advice on the AI models or parameters being used ⁵⁵ – for instance, if a certain model's output seems off or too slow, this agent can recommend switching models or adjusting settings. It logs its recommendations (MODEL_ADVICE) but does not enforce them automatically ⁵⁵.
- The *Troubleshooting Agent* is essentially the **incident response unit**. If the Orchestrator encounters a severe issue (an agent failing repeatedly, or a safety concern, etc.), it calls the Troubleshooting Agent, which generates an **INCIDENT_REPORT** describing the problem (symptoms, likely causes, suggested actions) ⁶². This report is then used to decide whether to escalate to a human. The Troubleshooting Agent does not fix anything on its own; it documents the incident and possibly stops the pipeline from proceeding ³⁷. This way, any failure is captured as a structured artifact for later analysis.
- Additionally, the architecture envisions roles like *Repository Archivist/Documentation Admin* (to manage repository state, ensure all changes are properly saved and documented) and possibly a *Handover Agent* (as hinted by an Orchestrator extension ⁶³, to prepare human-friendly summaries or next-session handover notes). These ensure that when the automated part of a workflow ends (or needs to pause), there is a clear record or package for the human team to review.

All these governance agents reinforce the idea that the system has an "**organized doubt**" principle: it systematically questions and records its own decisions, much like a built-in "red team" ³⁰. The Red-Team

concept in the strategy is reflected by the Challenger agent and governance reviews; while not a separate agent per se, the process encourages challenging assumptions and double-checking via these mechanisms.

- **Data/Information Layer:** Underpinning the agents is the management of content and data. The system operates on several key data artifacts:
 - The **Content Repository**: The digital text of Mustika Rasa across its versions. Initially, this includes the raw OCR text and images of the original (Phase 0 data). As Phase 0 is completed, a *master digital text* (corrected OCR, structured properly) becomes the baseline ⁸. Phase 1 adds layers of annotation to produce an annotated text, and Phase 2 produces the final public text. The architecture stores each of these versions, or at least the differences, in a repository (likely a Git-based repository given the references to docs and versioning). Each recipe or section can be an identifiable unit (with an ID) that the system processes and tracks through phases. The **handover between phases** is explicit: e.g., Phase 1 output becomes input to Phase 2, etc. ⁶⁴. This ensures continuity and that no context is lost between stages.
 - **Glossary and Terminology Database**: A collection of terms (ingredient names, techniques, etc.) with their translations or standard forms. The Glossary Agent interacts with this – it might read current preferred terms and propose new ones. Importantly, changes to this database are managed via lifecycle (they might remain in a *proposed* state until approved) ²⁵. This database ensures consistency of terminology across the project.
 - **Reference Corpus**: The historical and research documents that provide context (for example, previous translations, academic papers on certain dishes, or other cookbooks). The Research Agent queries this corpus to produce its reports ⁴⁵. The architecture likely includes a directory or database of these reference files, and a mechanism for the Research Agent to search and cite them. Each `RESEARCH_REPORT` might include citations to specific files or lines (evidence), which then become part of the methodology log or annotation footnotes ⁴⁶.
 - **Audit Logs and Governance Records**: All outputs from governance agents (incident reports, methodology logs, model advice, human review logs) are stored as part of the project documentation. For example, the Methodology Archivist might append entries to a `decision_log.md`, and the Troubleshooting Agent might create timestamped incident files. These reside in the repository's governance folder (the prompt references `docs/10-governance/*` for policies and presumably logs).
 - **Prompt Definitions**: Interestingly, even the prompts for agents are treated as data in the system. All agent instructions (goals, backstories, output format expectations) are stored in markdown files in a `prompts/` directory and considered the *source of truth* for agent behavior ⁶⁵ ⁶⁰. The application code loads these prompts at runtime, meaning the behavior of an agent can be updated by editing a prompt file rather than changing code. This not only makes iteration easier, but is a governance feature: prompt changes can be reviewed and version-controlled. The architecture ensures no hidden prompts are hard-coded – reinforcing transparency in how AI decisions are made ³⁸.
 - **Execution Metadata**: Each run of the system (say processing a batch of recipes) is itself documented. There are references to run IDs, CLI commands used, configuration files (`.yaml` configs for models) ⁵⁴, etc. The architecture likely logs each run's configuration and outcome in a session log (indeed a `CODEX_SESSION_LOG.md` is maintained with chronological entries of actions taken ⁵). This is important for reproducibility – anyone should be able to see what version of the prompts and models were used for a given output.

The data architecture centers on **high traceability** and **version control**. The use of markdown files for both content and logs suggests that the entire project is maintained in a repository, where each change (to content or to the process) is tracked. This aligns with the vision of an auditable knowledge infrastructure. It

also means the system can produce **artifacts for each stage** that are human-readable (e.g., the annotated text as a markdown or XML, the logs as markdown), enabling easy review and even publication of the process alongside the product.

- **Technical Infrastructure Layer:** Finally, at the base, we have the technology that runs all of this. The Mustikarasa system is implemented in **Python**, with a custom framework (`mustikarasa_agents.py` and related scripts) to instantiate agents and run the orchestrated workflow ⁶⁰. The AI agents are powered by **Large Language Models (LLMs)**. The architecture is model-agnostic to some extent: it can work with different LLM backends (OpenAI GPT models, or local models like Mistral via the CrewAI/Ollama setup, as indicated in pilot documents). The *Technical Advisor* agent's role underlines that the model can be switched – e.g., choosing a smaller model for certain tasks for cost, or a more capable one when needed ⁵⁵ – so the system is designed to plug in different AI models without changing the agent logic, by abstracting model calls behind each agent.

In practice: - Each agent likely runs by sending its prompt (goal/backstory plus current task data) to an LLM and getting a response. - The Orchestrator coordinates these calls and passes data between agents. There might be an in-memory structure or simply files that are written and read (for instance, one agent's output becomes the input file for the next). - The environment might be a local runner that can also call out to cloud APIs if needed. The mention of *sandbox* and *crew runner* suggests a containerized or isolated environment for test runs ⁶⁶ ⁵⁴, which is a good practice for safety: they test new agent behaviors in a sandbox with no external side effects (no file writes, no tool calls, just JSON outputs) ⁶⁶. This indicates a **staging vs production** consideration: new or experimental capabilities (like a new annotation model) are first run in a constrained setting (Phase 3 pilot) before being integrated into the main workflow. The architecture thus supports different execution modes (pilot sandbox vs production run) to ensure stability. - The **infrastructure** could be a single machine or cloud VM running the Python code. Given the need for traceability, likely it's integrated with a version control system (Git) for all content and logs. There may also be continuous integration hooks for updating documentation when code changes (the session log suggests some automation in updating docs). - For storage, because everything is text (markdown, JSON, etc.), a file-based repository suffices. Large assets like images are also referenced (the system links images of pages/illustrations – these would be stored in a media folder and referenced by the text). - **Integration points:** The system might integrate with human users through the repository or command-line. For example, when a glossary decision needs human input, a file might be created (a glossary proposal file) which a human editor would then edit or approve. Or a notification could be sent to a project management interface – but from the docs, it appears many interactions are document-driven (the human gate logs, the GO/NO-GO forms in Phase-4 references ⁶⁷ are markdown files to be filled out by humans). This is a design choice to keep even human interactions within the auditable system, rather than offline.

In summary, the architecture can be visualized as a **pipeline of AI agents** orchestrated by a central controller (Orchestrator), surrounded by a rigorous governance scaffold (rules, logs, human checks). Data flows through this pipeline in stages (original text → restored text → annotated text → public text), and at each stage, the transformations are logged and justified. The modular design (each agent is a module) provides a strong *architectural runway*: new agents or updated prompts can be incorporated as the project evolves without disrupting the overall structure, because the Orchestrator and governance framework handle them in a uniform way. This fulfills the long-term vision of having not just a one-time result, but a sustainable system that can adapt and be reused ⁶.

Capabilities and Maturity

The Mustikarasa system encompasses a range of **capabilities**, each corresponding to specific functionalities often implemented as agents or modules. Below, we list the key capabilities along with their current maturity level in the project (as of the latest known status). The maturity is categorized as **Active (Implemented)**, **Designed (Planned but not fully implemented)**, or **Conceptual (Identified for future)** ⁴⁹. This gives a picture of which parts of the architecture are operational now and which are on the roadmap:

- **Active (Implemented) Capabilities:** *Maturity: Active – These capabilities have been developed and integrated into the current system (agents exist with working code and prompts)* ⁶⁸.
- **Orchestrator (Workflow Management):** *Maturity: Active – The Orchestrator agent is fully implemented and coordinates the end-to-end editorial workflow* ⁶⁰ ⁶⁹. It enforces template structure, logs steps, and triggers governance as needed, effectively acting as the central brain of the system.
- **Translation Quality:** *Maturity: Active – The Translation Quality Agent ensures the initial translation from Indonesian (with English references) to Dutch is accurate. It's implemented and actively evaluates or refines translations for fidelity to the source meaning.*
- **Readability Editing:** *Maturity: Active – The Readability Editor agent is implemented to polish the text into fluent, modern Dutch* ⁶⁰. It improves phrasing and clarity while respecting content (no unauthorized changes to meaning).
- **Cultural-Historical Annotation:** *Maturity: Active – Sometimes called Cultural-Historical Editor, this agent is active and adds annotations regarding cultural context or historical notes to the text. It ensures the edition is enriched with relevant background info for deeper understanding.*
- **Structural Formatting (Book Structure):** *Maturity: Active – The Book Structure Agent is active, handling structural elements of the text (chapter/section demarcations, formatting consistency, table of contents alignment, etc.). It may ensure that each recipe follows the template sections (ingredients, steps, notes) and that overall book structure is intact.*
- **Recipe-Specific Editing:** *Maturity: Active – The Recipe Editor agent focuses on the recipe content itself (e.g., ensuring ingredients and instructions are clear and perhaps standardized). It's implemented to handle nuances of recipe text.*
- **Table Handling:** *Maturity: Active – The Table Editor agent is active, responsible for reconstructing or formatting tables from the original (like nutritional tables or ingredient tables)* ⁷⁰. It ensures tables in the book are captured correctly and legibly in the digital format.
- **Image Integration:** *Maturity: Active – The Image Agent is implemented to manage images (linking scanned images or illustrations to the text, verifying image references)* ⁷⁰. It keeps track of figures and ensures that where the book has images, the digital system associates them properly.
- **Continuity/Cohesion Checking:** *Maturity: Active – The Continuity/Cohesion Agent (the role sometimes split as Continuity Agent and Cohesion Agent) is active in code* ⁷¹ ⁴⁹. It checks for internal consistency, such as consistent use of names, units, and narrative tone across the entire content, and ensures that the final compilation feels coherent as a single edition.
- **Fidelity Assurance:** *Maturity: Active – The Fidelity Agent is active and compares the final Dutch text against the original (and possibly the English interim text) to flag any meaning shifts* ⁵⁸ ⁵⁹. It produces an "Opmerkingen" (Notes) section with potential meaning deviations ⁵⁸, thereby assuring that the modernization hasn't introduced mistranslations or distortions.
- **Design and Layout:** *Maturity: Active – The Design Agent is active, focusing on stylistic and layout considerations. For instance, it may ensure that the output text adheres to a style guide or that it's*

formatted for the intended medium (print or digital). This agent helps bridge content with presentation (important for the coffee-table quality of the public edition ¹⁶).

- **Critical Challenger:** *Maturity:* Active – The Challenger Agent is active and serves as an internal critic or double-check. It might re-evaluate translations or edits from a different angle to catch issues the primary pass missed, or pose questions ("Are we sure about X?") that force either the AI or a human to review certain decisions. Its prompt is implemented, and it's part of ensuring "*organized doubt*" in the system.
- **Designed (Prompt Available, Integration Pending) Capabilities:** *These capabilities have been defined (with prompt and intended function) but are not yet fully integrated into the automated workflow (they may be in testing or awaiting development)* ⁷².
- **Methodology Archivist (Process Logger):** *Maturity:* Designed – The prompt for the Methodology Archivist exists (to output a structured METHODOLOGY_LOG of decisions and context) ³⁴, but full integration is pending. This agent will automatically document important governance events and rationale during runs. Currently, some logging is done by the Orchestrator directly, but the dedicated agent will enhance and formalize the audit trail.
- **Technical Advisor (Model Strategy):** *Maturity:* Designed – A Technical Advisor agent has been designed to output MODEL_ADVICE on model choices and parameters ⁷³. It's not yet active in the live pipeline. When integrated, it will monitor for situations like a model underperforming or needing a different approach (e.g., suggesting to use a bigger model for translation vs a smaller one for routine edits) and will log recommendations ⁵⁵.
- **Troubleshooting Agent (Incident Responder):** *Maturity:* Designed – The Troubleshooting Agent's role and output format (INCIDENT_REPORT) are defined ⁷⁴, and the Orchestrator is prepared to call it on errors ⁷⁵ ³⁷. It has been tested in pilot scenarios, but integrating it fully in production runs is in progress. Once active, it will handle all exceptions by logging them and preventing uncontrolled continuation.
- **Glossary / Terminology Manager:** *Maturity:* Designed – The Glossary Agent (Terminology Manager) is designed to propose glossary updates and ensure consistent term usage ⁷⁶. Its prompt format (GLOSSARY_PROPOSALS) is set ⁷⁶. Full integration is pending completion of the glossary lifecycle governance; currently, glossary changes are likely handled in a semi-manual pilot mode ⁷⁷. Bringing this agent fully online will enable automated suggestions for term standardization, with human approval steps enforced.
- **Research / Historical Knowledge Agent:** *Maturity:* Designed – The Research Agent's specifications are in place (it provides a RESEARCH_REPORT with scope, summary, relevant files, insights, etc.) ⁷⁸. It has been tested in isolated cases (e.g., a pilot to see how research and glossary can work together ⁷⁷), but is not yet routinely called in every run. When integrated, it will fetch and supply contextual information for annotations and decisions, greatly enriching the system's outputs with evidence.
- **Repository Archivist / Documentation Admin:** *Maturity:* Designed – This role is outlined to handle tasks like finalizing documentation, managing the repository state after runs (committing changes, updating indexes). While less is mentioned about its implementation, it's listed as designed (prompt present) but not active ⁷⁹. It will be important when moving from pilot mode to continuous operations, ensuring that each run's results are properly saved and that any documentation (like a change log or readme) is updated accordingly.

- **Conceptual (Future/Planned) Capabilities:** These are identified needs or ideas for which detailed design may not yet exist. They represent the future extensions of the system ⁸⁰.
- **OCR / Source Integrity Agent:** Maturity: Concept – This capability would automate Phase 0 (source restoration) tasks, such as scanning quality control and OCR error detection. It's conceptual because early phases likely did a lot of OCR correction manually or with one-off scripts. A dedicated agent could, in the future, intelligently detect text recognition errors or missing elements by cross-referencing image scans. Ensuring **source integrity** (no loss or corruption of original content) is crucial, and an AI agent could assist in that, but designing it is a future work item.
- **Batch Governor:** Maturity: Concept – As the project scales to full-book operations, a Batch Governor agent is envisioned to manage the processing of batches of recipes or chapters. This agent would orchestrate at a higher level than the Orchestrator, perhaps launching multiple orchestrator instances or scheduling tasks, and handling batch-level concerns like overall progress, combined outputs, and partial failures. It's currently a concept, with the need recognized from scenario tests (like ensuring one failing recipe doesn't derail an entire batch) ⁵². Implementing this will improve throughput and reliability when moving from single-document runs to processing the whole book in an automated way.
- **Meeting / Redaction Agent:** Maturity: Concept – This is an interesting future concept likely tied to **human collaboration interfaces**. The "Meeting Agent" could facilitate structured discussions or reviews (for instance, summarizing a review meeting's conclusions, or orchestrating a conversation between multiple experts and the system). The "Redaction" part (in Dutch, *redactie* means editorial) suggests assisting in final editorial decisions or compilation. Possibly, this agent would simulate or support an editorial board meeting: aggregating different agent reports and human comments into a coherent action list or final changes. This is still an idea and not concrete, but underlines a long-term vision where AI could help coordinate between multiple humans (editors, translators, subject experts) in making final content decisions.

This breakdown shows that the **core content processing capabilities are largely in place (many Active)**, whereas **governance and scaling capabilities are being built out** (Designed or Concept). The project has consciously started with the core pipeline (translation through fidelity) and basic governance, and is now adding the more advanced governance (logging, research, glossary lifecycle) and scaling (batch processing, collaborative editing) features. The **maturity levels** highlight the *architectural runway*: the system's foundation anticipates those future capabilities, and steps are being taken to integrate them when ready ⁴⁹. For instance, prompts for agents like Methodology Archivist and Research Agent exist so that they can be plugged in with minimal friction once tested. This phased maturation is deliberate, ensuring stability at each stage.

Future Roadmap and Architectural Runway

The Mustikarasa architecture is built to support not only the immediate phases of the project but also future growth and adaptation. The concept of an "**architectural runway**" refers to preparing enough foundational architecture so that upcoming requirements can be implemented without major redesign. In this project, the runway is evident in how the architecture aligns with the planned phases and beyond:

- **Phase 3 (Pilot and Refinement):** According to the project strategy, after establishing the core workflow (Phases 0-2), **Phase 3** involves controlled pilots to test the governance and new agents in practice ⁸¹. The architecture is currently in this stage: for example, a pilot was conducted for

glossary and research integration, demonstrating the lifecycle without making permanent changes ⁷⁷. The system's design supports such pilots by having toggleable or isolated components (e.g., running the Research Agent in a read-only mode to see its output without affecting the main workflow). A *Phase-3 framing document* has been established as a strategic baseline ⁸², which the architecture adheres to by imposing boundaries (like all pilot changes are documentary only, not altering production data). This ensures that the outcomes of Phase 3 (lessons learned, refined rules) can smoothly transition into the main system.

- **Phase 4 and 5 (Integration and Scale-Up):** Although not explicitly named in the vision, internal documents suggest a Phase 4 where the system begins to handle real runs at scale (perhaps a sandbox full run, then production runs), and Phase 5/6 where the process is consolidated and scaled out. The **architectural runway** includes:
 - Implementing the **governance agents** in the runtime (Methodology Archivist, Troubleshooting, Technical Advisor) so that by Phase 4 the system can run with full audit and safety nets, even if humans are less hands-on per run.
 - **Scaling to larger batches:** The Batch Governor concept will likely come into play. The architecture already isolates recipe-level processing, which means in Phase 4+ it's feasible to run multiple recipes in parallel or in a sequence without manual intervention, as long as logging and stop mechanisms handle issues. Indeed, a *Phase-6 Status* note indicates the focus is on a governed, excerpt-aware editorial workflow, with traceability and human review prioritized over raw speed ⁸³. This confirms that the architecture is moving towards handling the entire book in a repeatable, safe way.
 - **Performance tuning:** As more agents (like Research) become integrated, the runtime will slow down if not optimized. The runway includes the ability to plug in optimized models (with Technical Advisor guidance) and possibly infrastructure scaling (e.g., running heavy analyses offline or using asynchronous calls). Since the architecture cleanly separates concerns (one agent = one function), components can be individually optimized (like swapping out the translation model for a faster one, or caching research results) without breaking the pipeline.
- **Phase 2 Public Edition and Publication:** Eventually, the system should be capable of producing the polished public edition automatically (with human approvals at key points). The architecture is already geared for this by having the Design Agent and readability/fidelity checks in place. What remains on the roadmap is final **productization**:
 - Generating the final publishable formats (e.g., PDF or web format of the book) – this might involve additional tooling outside the core AI agents, like formatting the combined output, embedding images properly, etc. The architecture's data layer (which keeps content in structured text and images) will feed into these publishing tools.
 - A possible **User Interface** for the public to consume the digital edition (or for editors to review content). While currently much of the system is command-line and repository-driven, a future step might be to create a web interface where the traceability ("how it was → how it is → why it changed") can be displayed to end users interactively. The architecture's careful logging and data structuring ensure that such a UI can be built on top of the existing data (for example, showing side-by-side original and edited text with annotations, drawn from the methodology logs and versioned content).
- **Knowledge Infrastructure extension:** After Mustika Rasa, the tools and architecture could be applied to other documents or projects. The runway considers this by avoiding hard-coding anything

specific to this cookbook's content. New agents or prompts could be written for a different book's unique needs, but the orchestrator-governance pattern would remain applicable. The long-term vision explicitly states the project is building an infrastructure, not just one translation ²², and the architecture's modular, document-driven approach is key to fulfilling that — it can onboard new content, new corpora, or new output requirements with minimal structural change.

- **Continuous Governance and Improvement:** As part of the roadmap, the team will continue to refine governance rules. The architecture is flexible enough to incorporate new rules – for instance, if a new kind of risk is identified, they could add a check either in an existing agent or by introducing another governance agent. The existence of a Red-Team review process as “organized doubt” ³⁰ might evolve into additional automated checks or more formalized human review steps. The **runway** includes the change management processes (like the requirements change process and phase framing documents) ²⁴ ⁸⁴ which ensure that as the architecture changes or new capabilities are turned on, everything remains aligned with the vision and properly documented. This meta-architecture governance means the architecture itself is under version control and review, preventing drift from the intended goals.

In essence, the architecture has been built with future phases in mind. It is *broadly extensible*: new agents can be integrated by adding prompt files and minimal code, thanks to the orchestrator's generalized handling of agents. It is *scalable*: by conceptualizing batch management and parallel workflows ahead of time. And it is *sustainable*: heavy emphasis on documentation, logging, and modular design ensures that the system can onboard new team members or be handed off, carrying its knowledge with it.

This **forward-looking design** gives confidence that as the Mustikarasa project moves into its next stages – completing the scholarly edition, releasing the public edition, and possibly tackling new projects – the architecture will support those needs without requiring a ground-up rebuild. The runway that has been built will allow the team to accelerate safely when ready, much like an airplane taking off once all checks are green.

Recommendations

Based on the current architecture and capability maturity, here are key recommendations to ensure the project's continued success and alignment with its goals:

- **1. Complete Integration of Governance Agents:** Expedite the implementation of the **Methodology Archivist, Troubleshooting Agent, Technical Advisor, Research Agent, and Glossary Agent** into the live workflow. These designed agents will greatly enhance traceability and reliability (e.g. automated decision logging, immediate incident reports, contextual knowledge injection). In particular, integrating the Methodology Archivist will ensure every significant decision or deviation is captured in the audit trail during runs (not just afterward), and the Troubleshooting Agent will enforce the ASR-006 rule that any incident triggers a documented stop ²³. Before moving to full-scale runs, ensure these agents are tested and their outputs are reviewed to fine-tune their prompts.
- **2. Implement the OCR/Source Integrity Capabilities for Phase 0:** To fully realize Phase 0 (Source Restoration), develop the OCR/Source Integrity Agent or equivalent tooling. This could involve training or prompting an agent to spot common OCR errors or layout issues by comparing scans to text, ensuring the *digital master* is as accurate as possible. Even if not fully automated, an agent that

flags suspicious OCR segments for human review would be valuable. This reduces the manual burden on editors and ensures no piece of text is left unverified. Since Phase 0 forms the foundation for all subsequent work, investing here pays off in fewer downstream content errors.

- **3. Develop the Batch Governor for Scaling:** As the project moves from handling individual recipes to processing the entire book, implement the **Batch Governor** concept to manage multi-recipe workflows. This component should handle task scheduling, parallelism (where safe), and error isolation across recipes. It can work in tandem with the Orchestrator: for example, the Batch Governor could launch one Orchestrator instance per chapter or per a set of recipes and monitor their completion. It should utilize the partial failure handling design ⁵² – meaning if one instance encounters a governance-stop, the Batch Governor can pause that instance and continue others, or collect all incidents to present for review at the end. A robust batch manager will help achieve efficiency without sacrificing the careful governance (the system can then potentially process many recipes overnight, pausing only those that need attention). This is crucial for meeting project timelines once you push towards publishing the public edition.
- **4. Enhance Human Review Interface and Workflow:** Strengthen the **Human Gate review process** by developing clearer interfaces or protocols for human interventions. Currently, human approvals and inputs seem to be managed through log files and documents (e.g., GO/NOGO files, HUMAN_GATE_LOG.md entries) ⁶⁷. It is recommended to create a more user-friendly dashboard or checklist for reviewers to use when the system stops for approval. For example, when a glossary term needs confirmation or an incident occurs, the system could generate a concise report (already happening via INCIDENT_REPORT and methodology logs) and present it in a dashboard where a human can make a decision (approve term, re-run step, etc.) which then feeds back into the system (perhaps by writing to an “approval” file that the Orchestrator reads). While this might be outside pure architecture and more in UI/UX, it’s important for efficiency: the easier it is for humans to interface with the system’s outputs and provide input, the faster the iterative cycles will go. Ensuring that this interface also logs decisions (which it should, by design) will maintain traceability.
- **5. Continue Pilot Testing and Knowledge Accumulation:** Use **Phase 3 pilots** to the fullest to refine prompts, policies, and agent behaviors. The current approach of running targeted mini-experiments (like testing how ambiguity is flagged, how multiple reviewers might conflict, how heavy annotation affects readability) ⁸⁵ ⁸⁶ is excellent. Keep designing such pilots for any open questions in the workflow (for instance, a pilot on how the Red-Team challenge could be formalized: simulate a scenario where the Challenger agent strongly disagrees with the Readability agent and see how the system handles it). Document the outcomes and adjust the architecture accordingly – e.g., if pilots show that certain decisions always require human input, maybe that step should be formalized as a governance-stop in the design. Pilots are a safe space to discover failure modes and corner cases; the recommendation is to address those findings in the architecture *before* full-scale execution. For example, if a pilot reveals a gap in the continuity check for cross-chapter references, one might introduce a new rule or a tweak in the Continuity Agent before doing the whole book.
- **6. Prepare for the Public Edition Output Format:** As Phase 2 (public edition) content is being generated, consider the final publication format requirements. Begin planning how the system’s output (which might be in Markdown/JSON now) will be transformed into a publishable form (PDF book, web site, etc.). This might involve collaboration with a design team for layout, but there are architectural considerations:

- Ensure that all content needed for publication is captured structurally (e.g., if sidebars or thematic introductions are planned, the system might need to accommodate creating those).
- Possibly develop an export tool or use an existing one to convert the annotated content into the desired format while preserving the traceability links (for instance, footnote numbering, appendices that list original text differences, etc.).
- Maintain the separation of concerns: the core system produces a content+metadata representation; a publishing module could take that and style it. This way, any changes in style do not require rerunning the AI pipeline, and vice versa. By anticipating these needs now, the team can ensure the data coming out of the AI workflow will be adequate for a high-quality public-facing product without last-minute rework.
- **7. Monitor and Refine the AI Models (Technical Advisor Feedback):** Leverage the Technical Advisor (once integrated) to keep the AI stack optimal. Over time, models improve; new versions of LLMs or domain-specific models might become available that could enhance quality or speed. The architecture is flexible enough to swap models (e.g., using a local model for drafts but a high-accuracy model for final verification). It's recommended to set up a periodic review (perhaps automated) of model performance using the logs: for instance, if the Challenger agent often catches issues from the Translation agent, that's an indicator to improve the translation model or prompt. Use the system's own data (methodology logs, incident reports) as feedback to pinpoint where the process might be bottlenecking or producing subpar results, and address those either by model improvements, prompt tuning, or agent logic adjustments. Essentially, adopt an **MLOps mindset**: monitor the outputs for quality and adjust the configuration to continuously improve.
- **8. Keep Strengthening Governance and Documentation:** As the system evolves, ensure to **update the Vision, Strategy, and Requirements documents** in parallel. The architecture should always be traceable back to an approved requirement or design decision. The project already has a strong practice of logging changes (e.g., using SYSTEM_REQUIREMENTS_CHANGE_PROCESS and logging session updates) ²⁴ ⁵. Continue this rigor. For any major architectural change (like introducing the Batch Governor or new agent), follow the change process: write a design proposal, review it (even via a red-team mindset), and get it approved. This will prevent scope creep or mission drift. Also, maintain the **Definition of Done criteria** in daily practice – e.g., don't consider Phase 2 complete until you can demonstrate, with the system's own logs, that all decisions are accounted for and all known risks flagged ⁵³. In short, use the architecture's strengths in traceability to manage the project itself: treat the architecture document (and allied strategy docs) as living documents that guide the implementation.

Implementing these recommendations will ensure that the Mustikarasa system not only achieves its immediate goals but also sets a benchmark for similar projects. By solidifying governance, completing pending capabilities, and planning for scale and publication, the team will be well-equipped to deliver a high-quality result that stands up to scrutiny from both scholars and the general public.

Conclusion

The Mustikarasa system's architecture is a robust blend of **innovative AI workflows and classical governance principles**. It marries an *agentic AI approach* – with specialized AI agents handling complex editorial tasks – to an *enterprise-grade oversight framework* that guarantees traceability, quality, and safety.

This document has outlined how the architecture meets the business and technical requirements drawn from the project's vision: from preserving the authenticity of a historical text to enabling modern accessibility, all while documenting every change "*so that future generations understand not just what was done, but why*".

By structuring the architecture along TOGAF-like domains, we have shown the logical organization: the business goals drive the functions needed, which in turn dictate the system components, data flows, and technical infrastructure. The current capabilities demonstrate that the core system is already functional and delivering results in line with Phase 1 objectives (scholarly edition), with Phase 2 (public edition) on the horizon. Meanwhile, the planned capabilities and recommendations illuminate a clear path forward to finalize the project and even generalize the solution to a broader knowledge infrastructure.

In conclusion, the Mustikarasa architecture provides a **complete blueprint** for turning raw historical data into a vetted, enriched publication. It exemplifies how to harness powerful AI tools in a constrained, transparent way – using agents to do the heavy lifting, but always under watchful governance and with humans in ultimate control. This architecture not only serves the immediate project needs but also lays down a template for similar future projects where preserving truth and context is as important as the end product. The strategy of combining "*top architect*" vision with deep research into requirements has resulted in a solid, extensible design. With continued diligent execution, the Mustikarasa system will achieve its ambitious goals and stand as a model for digital humanities projects powered by AI.

1 2 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 22 26 27 28 29 30 31 53 70

VISION_AND_STRATEGY.md

file:///file_0000000270072469b2bf4baf7037cf

3 4 25 33 35 37 38 39 40 41 42 43 44 45 46 47 50 51 55 61 62 63 75 orchestrator.md

file:///file_00000006ea472439184422037b81714

5 23 24 32 36 52 56 77 81 82 84 CODEX_SESSION_LOG.md

file:///file_0000000774c71f4a4f7f0bbf1547201

21 83 P6_PHASE_STATUS_SNAPSHOT.md

file:///file_000000001b047243abb6cac252bbfa9d

34 48 49 58 59 60 64 65 68 69 71 72 73 74 76 78 79 80 total_project.md

file:///file_000000001fd871f4b4f93951c179cdad

54 P6_CASE01_REAL_RUN_EXECUTION_PLAN.md

file:///file_00000005ee071f4ba24d08f4a1f34ab

57 Orchestrator-Worker Agents: A Practical Comparison of Common Agent Frameworks - Arize AI

<https://arize.com/blog/orchestrator-worker-agents-a-practical-comparison-of-common-agent-frameworks/>

66 67 85 86 CODEX_TODO.md

file:///file_000000079d471f4bb3ed86cd8db7a77